

Take-Home Case Study

Conversational AI Operations Specialist

OVERVIEW

Fora uses an AI advisor support chatbot called Sidekick (built on Intercom's Fin) to help travel advisors with questions about processes, policies, and booking guidance. As Sidekick's capabilities expand, we need systematic approaches to ensure response quality, identify content gaps, and maintain accuracy at scale.

This exercise simulates a real operational challenge: building a QA framework to audit Sidekick's performance on a complex topic (group bookings), then applying it to identify and fix quality issues.

BACKGROUND

Group bookings are one of the more complex areas advisors ask about. Questions range from basic definitions ("What counts as a group?") to nuanced contract negotiations ("What attrition rate should I accept?") to cruise-specific terms ("How do TC credits work?").

Sidekick pulls from a knowledge base of articles, training materials, and policy docs—but not all content is equally accurate or up-to-date. Some articles conflict with each other. Some are missing important nuance. Some questions don't have clear answers in the available content.

Your task is to design a QA system to audit Sidekick responses, then apply it to identify patterns in where Sidekick succeeds and fails.

WHAT YOU'LL RECEIVE

1. Mock knowledge base (~10 articles about group bookings)

- Varying publication dates (some recent, some outdated)
- Varying quality (some accurate, some with errors or missing context)
- One article labeled "Source of Truth" (fully accurate, but don't assume it's comprehensive or covers all topics)

2. Sample conversations (50 advisor–Sidekick exchanges)

- Real questions advisors might ask
- Sidekick's responses based on the knowledge base
- Some responses are good, some have issues

3. Additional knowledge sources (4 items)

- 2 chat exchanges between advisors and experts

- 1 email thread (HQ + advisor)
- 1 Slack conversation
- These represent accurate, high-context information

YOUR DELIVERABLES

Deliverable 1 — QA Framework & Solution

PRIMARY DELIVERABLE — SUBMIT AS A SHAREABLE, WORKING LINK

Design an operational system to audit and improve Sidekick response quality. Submit this as a working tool we can open and use—a Google Sheet, Notion workspace, coded app, or similar. A description of how you'd build it is not sufficient.

Your system should:

- Identify problems: How do you systematically find inaccurate, outdated, or incomplete responses?
- Categorize issues: What types of errors exist? How should they be prioritized?
- Propose fixes: For each issue found, what's the fix? (Update article, delete content, add new source, flag as unanswerable, etc.)
- Support multi-team collaboration: Groups involve Tour Ops, Booking Platform, and Support teams. How does your framework accommodate input from multiple stakeholders?

You can use any tools or approaches you want: spreadsheets, rubrics, AI-assisted analysis, manual review processes, testing protocols, etc. Focus on building something repeatable and scalable, not just fixing individual responses.

Deliverable 2 — Issue Log + Updated Content Base

THREE COMPONENTS IN ONE SUBMISSION

Run your QA system on the 50 sample conversations and knowledge base.

- Issue log in the specified format: Conversation # | Issue description | Proposed fix
 - Example: "Conversation 13 | Wrong cruise minimum (says 5+ cabins, should be 8+ staterooms) | Update Article 2"
- Updated content base: Resubmit all 10 articles with your edits applied.
- Pattern analysis: What systemic issues emerged across the conversations? What are the top 3–5 patterns causing quality problems?

If something isn't certain, flag it for expert review or specify where you made assumptions.

Deliverable 3 — Technical Documentation

MAX 3 PAGES — WRITTEN AS A RUNBOOK, NOT A REFLECTION

Document how your QA framework works so another team member could operate it:

- How does the system work technically? (What tools, what inputs/outputs, what happens at each step?)
- How do you run it start to finish? (Step-by-step operational guide)

- Where does AI do the work vs. where do humans need to intervene?
- What are the key decision points or judgment calls?

HOW AI AND HUMANS WORK TOGETHER

A realistic QA system uses AI to handle volume and humans to handle judgment. For example:

- AI: Categorize conversations, flag potential issues, group errors, identify contradictions in the knowledge base, and re-run conversations after updates to detect changes in response quality.
- Humans: Review flagged responses, resolve ambiguous cases, determine which sources are correct, prioritize fixes, and calibrate evaluation criteria over time.

In production, the humans in the loop are typically subject matter experts—support specialists in specific queues, Tour Ops stakeholders, and Booking Platform team members. They're best equipped to verify accuracy and make judgment calls, but they're also busy. When your framework outputs recommendations for human review, think about making them digestible, easy to understand, and shareable across teams.

For this exercise, you're the human in the loop—use AI tools however you want, but demonstrate you can make good judgment calls on what AI surfaces.

IN-PERSON COMPONENT

If your submission advances to the final round, one portion will be a 30–45 minute working session where you'll apply your QA system to 10 new conversations in real time.

You'll have access to the knowledge base (including your updates), and we'll observe how you work through your process.

This portion simulates the actual job: using your framework to triage new issues as they come in. Some conversations will have clear answers in your updated KB, some won't—just like in production.

NOTES

- No perfect answers: Some questions in the sample conversations are ambiguous or unanswerable with the given content. Part of the exercise is recognizing when that's the case.
- Think in systems: We're less interested in whether you catch every individual error and more interested in whether you build a framework that could catch them at scale.
- Use your judgment: If you think the exercise is missing something important, call it out. If you'd approach this differently in a real-world scenario, explain why.